

# DATA\_598\_HW\_4

January 30, 2022

## 1

Homework 4: AutoEncoders

Apoorv Sharma

<center> DATA 598 (Winter 2022), University of Washington </center>

## 2 1. Denoising AutoEncoders and Step Decay Learning Rates

In this exercise, we will use autoencoders to denoise (= de-noise, or remove the noise from) an image. We will also implement a step decay learning rate, a commonly used trick (for all deep kinds of deep nets, not just autoencoders).

Suppose we have an image  $x$  and *corrupt* it by some means to get  $x' = C(x)$ . Example corruptions including adding Gaussian noise or deleting random patches in the image. A denoising autoencoder with encoder  $h_w$  and decoder  $g_v$  (with respective parameters  $w$  and  $v$ ) takes in the corrupted input  $x'$  and returns  $\hat{x} = g(h(x'))$  that approximates the noise-less image  $x$ .

We will train a denoising autoencoder to reconstruct the noiseless images from the noisy ones, by minimizing the corresponding reconstruction error:

$$\min_{w,v} \mathbb{E} \|x - g_v \circ h_w(C(x))\|^2$$

We will use a step decay learning rate schedule

$$\gamma_t = \frac{\gamma_0}{2^{\lfloor t/t_0 \rfloor}}$$

in epoch  $t$ , where  $\gamma_0$  is a given initial learning rate, and  $t_0$  threshold. The learning is cut by a factor of 2 every  $t_0$  epochs:

$$\gamma_0 \dots \gamma_0, \frac{\gamma_0}{2} \dots \frac{\gamma_0}{2}, \frac{\gamma_0}{4} \dots \frac{\gamma_0}{4}$$

A larger learning rate makes faster progress initially whereas a smaller learning rate is more helpful closer to convergence. The step-decay schedule aims to get the best of both worlds.

```
[1]: import torch
    from torch.nn.functional import relu
```

```

from torchvision.datasets import MNIST

import numpy as np
import matplotlib.pyplot as plt

import math
import pickle

```

### 2.0.1 Download and Process MINST dataset

Perform the same preprocessing as in this week's lab.

```

[2]: # download dataset (~117M in size)
train_dataset = MNIST('./data', train=True, download=True)
X_train = train_dataset.data # torch tensor of type uint8
y_train = train_dataset.targets # torch tensor of type Long
test_dataset = MNIST('./data', train=False, download=True)
X_test = test_dataset.data
y_test = test_dataset.targets

# choose a subsample of 10% of the data:
idxs_train = torch.from_numpy(
    np.random.choice(X_train.shape[0], replace=False, size=X_train.shape[0]//
↳10)).long()
X_train, y_train = X_train[idxs_train], y_train[idxs_train]
# idxs_test = torch.from_numpy(
#     np.random.choice(X_test.shape[0], replace=False, size=X_test.shape[0]//
↳10))
# X_test, y_test = X_test[idxs_test], y_test[idxs_test]

print(f'X_train.shape = {X_train.shape}')
print(f'n_train: {X_train.shape[0]}, n_test: {X_test.shape[0]}')
print(f'Image size: {X_train.shape[1:]}')

```

```


X_train.shape = torch.Size([6000, 28, 28])
n_train: 6000, n_test: 10000
Image size: torch.Size([28, 28])

```

```

[3]: f, ax = plt.subplots(1, 5, figsize=(20, 4))
for i, idx in enumerate(np.random.choice(X_train.shape[0], 5)):
    ax[i].imshow(X_train[idx], cmap='gray', vmin=0, vmax=255)
    ax[i].set_title(f'Label = {y_train[idx]}', fontsize=20)

```



```
[4]: # NOTE: Run this cell only once. If you have to rerun it multiple times,
# make sure that you run it after rerunning the previous 2 cells.

# Normalize dataset: pixel values lie between 0 and 255
# Normalize them so the pixelwise mean is zero and standard deviation is 1

X_train = X_train.float() # convert to float32
# NOTE: we are returning a single mean/std over all the pixels, rather than a
# pixel-wise one
mean, std = X_train.mean(), X_train.std()
X_train = (X_train - mean) / (std + 1e-6) # avoid divide by zero
# X_train /= torch.norm(X_train, dim=(1, 2)).max()

X_test = X_test.float()
X_test = (X_test - mean) / (std + 1e-6)
# X_test /= torch.norm(X_test, dim=(1, 2)).max()

n_class = np.unique(y_train).shape[0]
```

## 2.0.2 Create convolutional autoencoder

Use the same convolutional autoencoder as in this week's lab, with a lower latent dimension of 40.

```
[5]: class EncoderModule(torch.nn.Module):
    def __init__(self, lower_dimension):
        super().__init__()
        # (B, 1, 28, 28) -> (B, 4, 12, 12)
        self.conv1 = torch.nn.Conv2d(1, 4, kernel_size=5, stride=2, padding=0)
        self.conv2 = torch.nn.Conv2d(4, 8, kernel_size=3, stride=2, padding=0)
        # Flatten (B, 8, 5, 5) -> (B, 8*5*5): do this in `forward()`
        # (B, 8*5*5) -> (B, lower_dimension); 8*5*5 = 200
        self.linear = torch.nn.Linear(200, lower_dimension)

    def forward(self, images):
        out = relu(self.conv1(images)) # conv1 + relu
        out = relu(self.conv2(out)) # conv2 + relu
        out = out.view(out.shape[0], -1) # flatten
```

```

out = self.linear(out)  # Linear
return out

```

```

[6]: class DecoderModule(torch.nn.Module):
    def __init__(self, lower_dimension):
        super().__init__()
        # (B, lower_dimension) -> (B, linear)
        self.linear_t = torch.nn.Linear(lower_dimension, 200)
        # Unflatten (B, 8*5*5) -> (B, 8, 5, 5); do this in `forward()`
        # (B, 8, 5, 5) -> (B, 4, 12, 12)
        self.conv2_t = torch.nn.ConvTranspose2d(8, 4, kernel_size=3, stride=2,
padding=0, output_padding=1)
        # (B, 4, 12, 12) -> (B, 1, 28, 28)
        self.conv1_t = torch.nn.ConvTranspose2d(4, 1, kernel_size=5, stride=2,
padding=0, output_padding=1)

    def forward(self, x):
        # Apply in reverse order
        out = relu(self.linear_t(x))  # linear_t + relu
        out = out.view(out.shape[0], 8, 5, 5)  # Unflatten
        out = relu(self.conv2_t(out))  # conv2_t + relu
        out = self.conv1_t(out)  # conv1_t (note: no relu at the end)
        return out

```

```

[7]: class AutoEncoder(torch.nn.Module):
    def __init__(self, lower_dimension):
        super().__init__()
        self.encoder = EncoderModule(lower_dimension)
        self.decoder = DecoderModule(lower_dimension)

    def forward(self, images):
        # Pass the images through the encoder to get the representations.
        # Then, pass the representations through the decoder to get the
reconstructed images
        # images -> encoder(.) -> decoder(.)
        out = self.encoder(images)
        out = self.decoder(out)

        return out

    def encode_images(self, images):
        return self.encoder(images)

    def decode_representations(self, representations):
        return self.decoder(representations)

```

### 2.0.3 Corruption function

As the corruption function  $C(\cdot)$ , we zero out a randomly chosen  $14 \times 14$  patch in the original image

```
[8]: def corrupt_image_batch(images):  
    # Add a 14x14 square of zeros in a 28x28 image  
    # images: (B, 1, 28, 28)  
  
    patch_size = 14 # zero out a 14x14 patch  
    batch_size = images.shape[0]  
    height, width = images.shape[-2:] # height and width of each image  
  
    starting_h = np.random.choice(height - patch_size, size=batch_size,  
    ↪replace=True)  
    starting_w = np.random.choice(width - patch_size, size=batch_size,  
    ↪replace=True)  
  
    images_corrupted = images.clone() # corrupt a copy so we do not lose the  
    ↪originals  
    for b in range(batch_size):  
        h = starting_h[b]  
        w = starting_w[b]  
        images_corrupted[b, 0, h:h+patch_size, w:w+patch_size] = 0 # set to 0  
    return images_corrupted
```

### 2.0.4 Training Functions

```
[9]: def loss_function(true_images, reconstructed_images): # square loss  
    residual = (true_images - reconstructed_images).view(-1) # flatten into a  
    ↪vector  
    # return the average over examples  
    return 0.5 * torch.norm(residual) ** 2 / (true_images.shape[0])  
  
def compute_objective(model, original_images, corrupted_images):  
    # reshape images from (B, 28, 28) -> (B, 1, 28, 28) as required by the model  
    reconstructed_images = model(corrupted_images)  
    return loss_function(original_images.unsqueeze(1), reconstructed_images)
```

```
[10]: @torch.no_grad()  
def compute_logs(model, verbose=False): # Only report loss  
    train_loss = compute_objective(model, X_train, corrupt_image_batch(X_train.  
    ↪unsqueeze(1)))  
    test_loss = compute_objective(model, X_test, corrupt_image_batch(X_test.  
    ↪unsqueeze(1)))  
    if verbose:  
        print('Train Loss = {:.3f}, Test Loss = {:.3f}'.format(  
            train_loss.item(), test_loss.item(),
```

```

    ))
    return (train_loss, test_loss)

def minibatch_sgd_one_pass(model, X, learning_rate, batch_size, verbose=False):
    num_examples = X.shape[0]
    average_loss = 0.0
    num_updates = int(round(num_examples / batch_size))
    for i in range(num_updates):
        idxs = np.random.choice(num_examples, size=(batch_size,))

        corrupted_images = corrupt_image_batch(X[idxs].unsqueeze(1))
        # compute the objective.
        objective = compute_objective(model, X[idxs], corrupted_images)
        average_loss = 0.99 * average_loss + 0.01 * objective.item()
        if verbose and (i+1) % 100 == 0:
            print("{:.3f}".format(average_loss))

        # Exercise:
        # compute the gradient using automatic differentiation
        gradients = torch.autograd.grad(outputs=objective, inputs=model.
↪parameters())

        # Perform the SGD update
        with torch.no_grad():
            for (w, g) in zip(model.parameters(), gradients):
                w -= learning_rate * g

    return model

```

### 2.0.5 Train Model

Train the model for 40 epochs starting with  $\gamma_0 = 2.5 \cdot 10^{-4}$  and  $t_0 = 10$  (i.e., halve the learning rate every 10 epochs).

```

[11]: initial_learning_rate = 2.5e-4
learning_rate_threshold = 10
batch_size = 1
lower_dimension = 40 # use a lower dimensionality of 40
num_epochs = 40

logs = []

model = AutoEncoder(lower_dimension)
print(f'Iteration 0, LR: {initial_learning_rate}', end=', ')
logs.append(compute_logs(model, verbose=True))

for j in range(num_epochs):

```

```

# step decay learning rate schedule
num_epoch = j + 1
learning_rate = initial_learning_rate / ( math.pow(2, math.floor(num_epoch/
↪learning_rate_threshold)) )

model = minibatch_sgd_one_pass(model, X_train, learning_rate, ↪
↪batch_size=batch_size, verbose=False)
print(f'Iteration {num_epoch}, LR: {learning_rate}', end=', ')
logs.append(compute_logs(model, verbose=True))

with open('./models/logs.pkl', 'wb') as f:
    pickle.dump(logs, f)

# save the model parms
torch.save(model.state_dict(), f'./models/parms.pt')

```

```

Iteration 0, LR: 0.00025, Train Loss = 403.148, Test Loss = 409.312
Iteration 1, LR: 0.00025, Train Loss = 105.786, Test Loss = 105.933
Iteration 2, LR: 0.00025, Train Loss = 98.256, Test Loss = 99.154
Iteration 3, LR: 0.00025, Train Loss = 97.623, Test Loss = 98.991
Iteration 4, LR: 0.00025, Train Loss = 86.478, Test Loss = 86.705
Iteration 5, LR: 0.00025, Train Loss = 85.741, Test Loss = 86.091
Iteration 6, LR: 0.00025, Train Loss = 89.762, Test Loss = 89.973
Iteration 7, LR: 0.00025, Train Loss = 86.330, Test Loss = 86.676
Iteration 8, LR: 0.00025, Train Loss = 78.222, Test Loss = 78.539
Iteration 9, LR: 0.00025, Train Loss = 71.088, Test Loss = 71.248
Iteration 10, LR: 0.000125, Train Loss = 79.126, Test Loss = 79.552
Iteration 11, LR: 0.000125, Train Loss = 74.888, Test Loss = 75.221
Iteration 12, LR: 0.000125, Train Loss = 68.541, Test Loss = 68.927
Iteration 13, LR: 0.000125, Train Loss = 69.009, Test Loss = 69.327
Iteration 14, LR: 0.000125, Train Loss = 72.473, Test Loss = 72.941
Iteration 15, LR: 0.000125, Train Loss = 66.059, Test Loss = 66.402
Iteration 16, LR: 0.000125, Train Loss = 67.637, Test Loss = 68.053
Iteration 17, LR: 0.000125, Train Loss = 68.926, Test Loss = 69.170
Iteration 18, LR: 0.000125, Train Loss = 67.436, Test Loss = 67.592
Iteration 19, LR: 0.000125, Train Loss = 69.775, Test Loss = 70.145
Iteration 20, LR: 6.25e-05, Train Loss = 64.511, Test Loss = 64.899
Iteration 21, LR: 6.25e-05, Train Loss = 66.767, Test Loss = 67.420
Iteration 22, LR: 6.25e-05, Train Loss = 67.294, Test Loss = 67.952
Iteration 23, LR: 6.25e-05, Train Loss = 70.986, Test Loss = 71.420
Iteration 24, LR: 6.25e-05, Train Loss = 65.656, Test Loss = 66.042
Iteration 25, LR: 6.25e-05, Train Loss = 66.357, Test Loss = 66.819
Iteration 26, LR: 6.25e-05, Train Loss = 61.469, Test Loss = 62.023
Iteration 27, LR: 6.25e-05, Train Loss = 63.393, Test Loss = 64.070
Iteration 28, LR: 6.25e-05, Train Loss = 64.305, Test Loss = 64.820
Iteration 29, LR: 6.25e-05, Train Loss = 63.250, Test Loss = 63.656
Iteration 30, LR: 3.125e-05, Train Loss = 63.937, Test Loss = 64.424

```

```
Iteration 31, LR: 3.125e-05, Train Loss = 65.244, Test Loss = 65.932
Iteration 32, LR: 3.125e-05, Train Loss = 63.428, Test Loss = 63.922
Iteration 33, LR: 3.125e-05, Train Loss = 63.397, Test Loss = 64.042
Iteration 34, LR: 3.125e-05, Train Loss = 62.125, Test Loss = 62.669
Iteration 35, LR: 3.125e-05, Train Loss = 65.279, Test Loss = 65.835
Iteration 36, LR: 3.125e-05, Train Loss = 62.458, Test Loss = 62.966
Iteration 37, LR: 3.125e-05, Train Loss = 62.087, Test Loss = 62.481
Iteration 38, LR: 3.125e-05, Train Loss = 63.737, Test Loss = 64.289
Iteration 39, LR: 3.125e-05, Train Loss = 63.072, Test Loss = 63.793
Iteration 40, LR: 1.5625e-05, Train Loss = 60.442, Test Loss = 60.926
```

View the loss


```
[12]: logs_arr = np.asarray(logs)

plt.plot(logs_arr[:, 0], label="Train")
plt.plot(logs_arr[:, 1], label="Test")
plt.title('AutoEncoder Loss', fontsize=18)
plt.ylabel('Loss', fontsize=18)
plt.xlabel('#Passes', fontsize=18)
plt.legend(fontsize=18)

# Try to set clever bounds
plt.ylim((logs_arr[1:].reshape(-1).min() * 0.95, logs_arr[1:].reshape(-1).max()
↪ * 0.95))
```

```
[12]: (57.420057106018064, 100.63650093078613)
```





## 2.0.6 Model Output

Show some examples of the denoising process from the test set.

```
[13]: f, ax = plt.subplots(3, 5, figsize=(20, 10))


idxs = np.random.choice(X_test.shape[0], 5)
images = X_test[idxs].unsqueeze(1)
images_corrupted = corrupt_image_batch(images)

for i, idx in enumerate(idxs):
    ax[0, i].imshow(images[i].squeeze() * std + mean, cmap='gray') # Note:
    ↳ Undo mean and std normalization before viewing image
    ax[0, i].set_title(f'Original #{idx}', fontsize=20)

    ax[1, i].imshow(images_corrupted[i].squeeze() * std + mean, cmap='gray') #
    ↳ Note: Undo mean and std normalization before viewing image
    ax[2, i].set_title(f'Corrupted #{idx}', fontsize=20)

    # add batch and channel dimensions before passing through the model and
    ↳ squeeze them out later
    xr = model(images_corrupted[i].view(1, 1, 28, 28)).detach().squeeze()
```

```
ax[2, i].imshow(xr * std + mean, cmap='gray') # Note: Undo mean and std,
↪normalization before viewing image
ax[2, i].set_title(f'Recons. #{idx}', fontsize=20)
```



## 3 2. (Bonus) AutoEncoders as a non-linear PCA

In this exercise, we will compare autoencoders versus PCA for dimensionality reduction. We will note their usefulness on the end goal of training a linear model using the extracted low-dimensional features.

In the first few labs, we used images as 784 dimensional vectors. Here, you will use either autoencoders or PCA on the training dataset to project the data onto a lower dimension  $d$ . You will then train a multinomial logistic regression model with scikit-learn and keep track of the test accuracy.

```
[14]: from sklearn.decomposition import PCA
      from sklearn.linear_model import LogisticRegression
```


```
[15]: num_components_list = [10, 25, 50, 100]
```

### 3.0.1 Download and Process MINST dataset

Perform the same preprocessing as in this week's lab.

```
[16]: f, ax = plt.subplots(1, 5, figsize=(20, 4))
      for i, idx in enumerate(np.random.choice(X_train.shape[0], 5)):
          ax[i].imshow(X_train[idx].squeeze() * std + mean, cmap='gray', vmin=0,
          ↪vmax=255)
```

```
ax[i].set_title(f'Label = {y_train[idx]}', fontsize=20)
```



### 3.0.2 PCA Projections

Given a lower dimension  $d$ , use PCA to reduce the dimensionality of the training set to  $d$  dimensions. Transform the test set by projecting on to the same space. You may use scikit-learn's PCA implementation.

```
[17]: X_train_2d = X_train.view((X_train.shape[0], -1))
      X_test_2d = X_test.view((X_test.shape[0], -1))
```

```
[18]: f, ax = plt.subplots(1, len(num_components_list)+1, figsize=(20, 10))
      idx = np.random.choice(X_train_2d.shape[0], 1)


      ax[0].imshow(X_train[idx].squeeze() * std + mean, cmap='gray') # Note: Undo
      ↪ mean and std normalization before viewing image
      ax[0].set_title(f'Num Components = {28*28}', fontsize=12)

      for i, num_components in enumerate(num_components_list):
          pca = PCA(n_components = num_components)
          principalComponents = pca.fit(X_train_2d)

          # Apply transform to both the training set and the test set.
          train_img = pca.transform(X_train_2d)
          test_img = pca.transform(X_test_2d)

          approximation = pca.inverse_transform(train_img)

          ax[i+1].set_title(f'Num Components = {num_components}', fontsize=12)
          ax[i+1].imshow(torch.tensor(approximation[idx].reshape(28,28)).squeeze()) *
          ↪ std + mean, cmap='gray') # Note: Undo mean and std normalization before
          ↪ viewing image
```



### 3.0.3 Multinomial logistic Regression

```
[19]: def get_accuracy(y_true, y_pred):
      return np.mean(y_pred == y_true)
```

```
[20]: log_reg_model = LogisticRegression(max_iter=200)
      log_reg_model.fit(train_img, y_train)
```

/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-packages/sklearn/linear\_model/\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[20]: LogisticRegression(max_iter=200)
```

```
[21]: y_pred = log_reg_model.predict(test_img)
      print(get_accuracy(y_pred, y_test.numpy()) * 100)
```

89.68

### 3.0.4 Define AutoEncoder and Training Functions

These need to be re-written/re-defined since they were changed for #1.

The same

```
[22]: def loss_function(true_images, reconstructed_images): # square loss
      residual = (true_images - reconstructed_images).view(-1) # flatten into a
      ↪ vector
      # return the average over examples
      return 0.5 * torch.norm(residual) ** 2 / (true_images.shape[0])
```

```
def compute_objective(model, images):
    # reshape images from (B, 28, 28) -> (B, 1, 28, 28) as required by the model
    images = images.unsqueeze(1) # Add channel dimension
    reconstructed_images = model(images)
    return loss_function(images, reconstructed_images)
```

```
[23]: @torch.no_grad()
def compute_logs(model, verbose=False): # Only report loss
    train_loss = compute_objective(model, X_train)
    test_loss = compute_objective(model, X_test)
    if verbose:
        print('Train Loss = {:.3f}, Test Loss = {:.3f}'.format(
            train_loss.item(), test_loss.item(),
        ))
    return (train_loss, test_loss)

def minibatch_sgd_one_pass(model, X, learning_rate, batch_size, verbose=False):
    num_examples = X.shape[0]
    average_loss = 0.0
    num_updates = int(round(num_examples / batch_size))
    for i in range(num_updates):
        idxs = np.random.choice(num_examples, size=(batch_size,))
        # compute the objective.
        objective = compute_objective(model, X[idxs])
        average_loss = 0.99 * average_loss + 0.01 * objective.item()
        if verbose and (i+1) % 100 == 0:
            print("{:.3f}".format(average_loss))

        # Exercise:
        # compute the gradient using automatic differentiation
        gradients = torch.autograd.grad(outputs=objective, inputs=model.
        ↪parameters())

        # Perform the SGD update
        with torch.no_grad():
            for (w, g) in zip(model.parameters(), gradients):
                w -= learning_rate * g

    return model
```

### 3.0.5 Dimension Reduction Analysis

Here we perform the following steps, for the following values of  $d$  [10, 25, 50, 100]: 1. Given a lower dimension  $d$ , use PCA to reduce the dimensionality of the training set to  $d$  dimensions. Transform the test set by projecting on to the same space. 2. Train an autoencoder using the same settings as the lab, but with a hidden dimension as  $d$ . Train it for 40 epochs. Use the encoder to obtain  $d$ -dimensional representations for all training and test images. 3. Train a multinomial

logistic regression model with scikit-learn using each of the representations you have obtained.

```
[28]: initial_learning_rate = 2.5e-4
      learning_rate_threshold = 10
      batch_size = 1
      num_epochs = 40
```

```
[29]: pca_accuracy = []
      ae_accuracy = []
```

```
[30]: for i, num_components in enumerate(num_components_list):

    print(f'Starting dimension reduction to {num_components} components')

    # STEP 1: Dimension Reduction using PCA
    print(f'\tStarting PCA')
    pca = PCA(n_components = num_components)
    principalComponents = pca.fit(X_train_2d)

    # Apply transform to both the training set and the test set.
    train_img_pca = pca.transform(X_train_2d)
    test_img_pca = pca.transform(X_test_2d)
    print(f'\tFinished PCA')

    # STEP 2: Dimension Reduction using AutoEncoder
    print(f'\tStarting AutoEncoder Training')
    logs = []
    model = AutoEncoder(num_components)
    print(f'\t\tIteration 0, LR: {initial_learning_rate}', end=', ')
    logs.append(compute_logs(model, verbose=True))

    for j in range(num_epochs):
        # step decay learning rate schedule
        num_epoch = j + 1
        learning_rate = initial_learning_rate / ( math.pow(2, math.
↪floor(num_epoch/learning_rate_threshold)) )

        model = minibatch_sgd_one_pass(model, X_train, learning_rate,
↪batch_size=batch_size, verbose=False)
        print(f'\t\tIteration {num_epoch}, LR: {learning_rate}', end=', ')
        logs.append(compute_logs(model, verbose=True))

    with open('./models/q2_logs_nc_{num_components}.pkl', 'wb') as f:
        pickle.dump(logs, f)

    # save the model parms
    torch.save(model.state_dict(), f'./models/q2_parms_nc_{num_components}.pt')
```

```

# Obtain encoded traina and test images
train_img_ae = model.encode_images(X_train.unsqueeze(1))
test_img_ae = model.encode_images(X_test.unsqueeze(1))

print(f'\tDone AutoEncoder Training')

# STEP 3a Create Logistic Regression Models for each reduction method
log_reg_model_pca = LogisticRegression(max_iter=200)
log_reg_model_pca.fit(train_img_pca, y_train)

log_reg_model_ae = LogisticRegression(max_iter=200)
log_reg_model_ae.fit(train_img_ae.detach().numpy(), y_train)

# STEP 3b: Predict and get accuracy from each model
# Accuracy for PCA model
y_pred = log_reg_model_pca.predict(test_img_pca)
pca_accuracy.append(get_accuracy(y_pred, y_test.numpy()) * 100)

# Accuracy for AutoEncoder model
y_pred = log_reg_model_ae.predict(test_img_ae.detach().numpy())
ae_accuracy.append(get_accuracy(y_pred, y_test.numpy()) * 100)

print(f'AE Acc: {ae_accuracy[-1]:.2f}, PCA Acc: {pca_accuracy[-1]:.2f}')

```

Starting dimension reduction to 10 components

Starting PCA

Finished PCA

Starting AutoEncoder Training

Iteration 0, LR: 0.00025, Train Loss = 398.408, Test Loss =	406.882
Iteration 1, LR: 0.00025, Train Loss = 388.082, Test Loss =	395.253
Iteration 2, LR: 0.00025, Train Loss = 125.771, Test Loss =	126.142
Iteration 3, LR: 0.00025, Train Loss = 118.647, Test Loss =	118.829
Iteration 4, LR: 0.00025, Train Loss = 121.273, Test Loss =	121.887
Iteration 5, LR: 0.00025, Train Loss = 114.497, Test Loss =	114.543
Iteration 6, LR: 0.00025, Train Loss = 112.021, Test Loss =	112.409
Iteration 7, LR: 0.00025, Train Loss = 110.072, Test Loss =	110.112
Iteration 8, LR: 0.00025, Train Loss = 108.094, Test Loss =	108.636

111.919	Iteration 9, LR: 0.00025, Train Loss = 110.723, Test Loss =
104.026	Iteration 10, LR: 0.000125, Train Loss = 103.235, Test Loss =
101.586	Iteration 11, LR: 0.000125, Train Loss = 101.062, Test Loss =
102.618	Iteration 12, LR: 0.000125, Train Loss = 101.835, Test Loss =
102.940	Iteration 13, LR: 0.000125, Train Loss = 101.978, Test Loss =
101.439	Iteration 14, LR: 0.000125, Train Loss = 100.311, Test Loss =
101.921	Iteration 15, LR: 0.000125, Train Loss = 100.948, Test Loss =
101.313	Iteration 16, LR: 0.000125, Train Loss = 100.228, Test Loss =
101.000	Iteration 17, LR: 0.000125, Train Loss = 100.060, Test Loss =
99.822	Iteration 18, LR: 0.000125, Train Loss = 99.171, Test Loss =
99.748	Iteration 19, LR: 0.000125, Train Loss = 98.942, Test Loss =
97.640	Iteration 20, LR: 6.25e-05, Train Loss = 96.578, Test Loss =
97.898	Iteration 21, LR: 6.25e-05, Train Loss = 96.684, Test Loss =
97.371	Iteration 22, LR: 6.25e-05, Train Loss = 96.339, Test Loss =
97.486	Iteration 23, LR: 6.25e-05, Train Loss = 96.418, Test Loss =
96.701	Iteration 24, LR: 6.25e-05, Train Loss = 95.683, Test Loss =
97.046	Iteration 25, LR: 6.25e-05, Train Loss = 95.896, Test Loss =
96.880	Iteration 26, LR: 6.25e-05, Train Loss = 95.799, Test Loss =
96.521	Iteration 27, LR: 6.25e-05, Train Loss = 95.466, Test Loss =
97.360	Iteration 28, LR: 6.25e-05, Train Loss = 96.302, Test Loss =
96.938	Iteration 29, LR: 6.25e-05, Train Loss = 95.937, Test Loss =
95.533	Iteration 30, LR: 3.125e-05, Train Loss = 94.265, Test Loss =
95.275	Iteration 31, LR: 3.125e-05, Train Loss = 94.139, Test Loss =
95.372	Iteration 32, LR: 3.125e-05, Train Loss = 94.069, Test Loss =



```

Iteration 33, LR: 3.125e-05, Train Loss = 94.028, Test Loss =
95.211
Iteration 34, LR: 3.125e-05, Train Loss = 94.090, Test Loss =
95.221
Iteration 35, LR: 3.125e-05, Train Loss = 93.887, Test Loss =
95.229
Iteration 36, LR: 3.125e-05, Train Loss = 93.858, Test Loss =
94.989
Iteration 37, LR: 3.125e-05, Train Loss = 93.658, Test Loss =
94.879
Iteration 38, LR: 3.125e-05, Train Loss = 93.673, Test Loss =
94.732
Iteration 39, LR: 3.125e-05, Train Loss = 93.820, Test Loss =
95.022
Iteration 40, LR: 1.5625e-05, Train Loss = 93.179, Test Loss =
94.494

```

Done AutoEncoder Training

```

/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(

```

AE Acc: 88.13, PCA Acc: 80.59

Starting dimension reduction to 25 components

Starting PCA

Finished PCA

Starting AutoEncoder Training

```

Iteration 0, LR: 0.00025, Train Loss = 392.124, Test Loss =
399.276

```

```

Iteration 1, LR: 0.00025, Train Loss = 110.442, Test Loss =

```

109.876	Iteration 2, LR: 0.00025, Train Loss = 77.722, Test Loss =
77.165	Iteration 3, LR: 0.00025, Train Loss = 73.069, Test Loss =
72.832	Iteration 4, LR: 0.00025, Train Loss = 72.282, Test Loss =
71.602	Iteration 5, LR: 0.00025, Train Loss = 70.758, Test Loss =
70.191	Iteration 6, LR: 0.00025, Train Loss = 65.048, Test Loss =
64.995	Iteration 7, LR: 0.00025, Train Loss = 62.681, Test Loss =
62.431	Iteration 8, LR: 0.00025, Train Loss = 61.486, Test Loss =
61.192	Iteration 9, LR: 0.00025, Train Loss = 61.099, Test Loss =
61.214	Iteration 10, LR: 0.000125, Train Loss = 57.536, Test Loss =
57.585	Iteration 11, LR: 0.000125, Train Loss = 56.843, Test Loss =
56.884	Iteration 12, LR: 0.000125, Train Loss = 56.359, Test Loss =
56.382	Iteration 13, LR: 0.000125, Train Loss = 56.539, Test Loss =
56.739	Iteration 14, LR: 0.000125, Train Loss = 56.082, Test Loss =
56.090	Iteration 15, LR: 0.000125, Train Loss = 55.765, Test Loss =
55.853	Iteration 16, LR: 0.000125, Train Loss = 55.421, Test Loss =
55.634	Iteration 17, LR: 0.000125, Train Loss = 55.756, Test Loss =
55.947	Iteration 18, LR: 0.000125, Train Loss = 55.017, Test Loss =
55.174	Iteration 19, LR: 0.000125, Train Loss = 55.691, Test Loss =
55.674	Iteration 20, LR: 6.25e-05, Train Loss = 53.242, Test Loss =
53.456	Iteration 21, LR: 6.25e-05, Train Loss = 53.222, Test Loss =
53.410	Iteration 22, LR: 6.25e-05, Train Loss = 53.244, Test Loss =
53.376	Iteration 23, LR: 6.25e-05, Train Loss = 53.491, Test Loss =
53.672	Iteration 24, LR: 6.25e-05, Train Loss = 52.973, Test Loss =
53.179	Iteration 25, LR: 6.25e-05, Train Loss = 53.386, Test Loss =

53.583  
 52.852  
 52.903  
 53.432  
 52.723  
 52.209  
 52.125  
 52.100  
 52.159  
 52.049  
 52.059  
 51.883  
 51.946  
 51.866  
 51.840  
 51.570

Iteration 26, LR: 6.25e-05, Train Loss = 52.634, Test Loss =  
 Iteration 27, LR: 6.25e-05, Train Loss = 52.682, Test Loss =  
 Iteration 28, LR: 6.25e-05, Train Loss = 53.113, Test Loss =  
 Iteration 29, LR: 6.25e-05, Train Loss = 52.535, Test Loss =  
 Iteration 30, LR: 3.125e-05, Train Loss = 51.927, Test Loss =  
 Iteration 31, LR: 3.125e-05, Train Loss = 51.827, Test Loss =  
 Iteration 32, LR: 3.125e-05, Train Loss = 51.801, Test Loss =  
 Iteration 33, LR: 3.125e-05, Train Loss = 51.868, Test Loss =  
 Iteration 34, LR: 3.125e-05, Train Loss = 51.747, Test Loss =  
 Iteration 35, LR: 3.125e-05, Train Loss = 51.760, Test Loss =  
 Iteration 36, LR: 3.125e-05, Train Loss = 51.637, Test Loss =  
 Iteration 37, LR: 3.125e-05, Train Loss = 51.615, Test Loss =  
 Iteration 38, LR: 3.125e-05, Train Loss = 51.505, Test Loss =  
 Iteration 39, LR: 3.125e-05, Train Loss = 51.488, Test Loss =  
 Iteration 40, LR: 1.5625e-05, Train Loss = 51.186, Test Loss =

Done AutoEncoder Training

/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-  
 packages/sklearn/linear\_model/\_logistic.py:814: ConvergenceWarning: lbfgs failed  
 to converge (status=1):  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-  
 regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

AE Acc: 90.73, PCA Acc: 88.52

Starting dimension reduction to 50 components

Starting PCA

Finished PCA

Starting AutoEncoder Training

400.636	Iteration 0, LR: 0.00025, Train Loss = 393.328, Test Loss =
76.322	Iteration 1, LR: 0.00025, Train Loss = 76.671, Test Loss =
62.395	Iteration 2, LR: 0.00025, Train Loss = 62.838, Test Loss =
56.715	Iteration 3, LR: 0.00025, Train Loss = 57.246, Test Loss =
54.016	Iteration 4, LR: 0.00025, Train Loss = 54.207, Test Loss =
75.448	Iteration 5, LR: 0.00025, Train Loss = 75.574, Test Loss =
49.592	Iteration 6, LR: 0.00025, Train Loss = 50.080, Test Loss =
54.251	Iteration 7, LR: 0.00025, Train Loss = 53.834, Test Loss =
46.436	Iteration 8, LR: 0.00025, Train Loss = 46.501, Test Loss =
46.062	Iteration 9, LR: 0.00025, Train Loss = 45.986, Test Loss =
39.671	Iteration 10, LR: 0.000125, Train Loss = 39.784, Test Loss =
38.953	Iteration 11, LR: 0.000125, Train Loss = 39.110, Test Loss =
38.656	Iteration 12, LR: 0.000125, Train Loss = 38.691, Test Loss =
38.016	Iteration 13, LR: 0.000125, Train Loss = 38.134, Test Loss =
38.085	Iteration 14, LR: 0.000125, Train Loss = 38.228, Test Loss =
37.737	Iteration 15, LR: 0.000125, Train Loss = 37.777, Test Loss =
37.096	Iteration 16, LR: 0.000125, Train Loss = 37.115, Test Loss =
	Iteration 17, LR: 0.000125, Train Loss = 36.782, Test Loss =

36.758	Iteration 18, LR: 0.000125, Train Loss = 36.945, Test Loss =
36.935	Iteration 19, LR: 0.000125, Train Loss = 36.630, Test Loss =
36.673	Iteration 20, LR: 6.25e-05, Train Loss = 35.518, Test Loss =
35.549	Iteration 21, LR: 6.25e-05, Train Loss = 35.370, Test Loss =
35.382	Iteration 22, LR: 6.25e-05, Train Loss = 35.105, Test Loss =
35.116	Iteration 23, LR: 6.25e-05, Train Loss = 35.148, Test Loss =
35.179	Iteration 24, LR: 6.25e-05, Train Loss = 34.953, Test Loss =
34.986	Iteration 25, LR: 6.25e-05, Train Loss = 34.871, Test Loss =
34.940	Iteration 26, LR: 6.25e-05, Train Loss = 34.837, Test Loss =
34.886	Iteration 27, LR: 6.25e-05, Train Loss = 34.756, Test Loss =
34.792	Iteration 28, LR: 6.25e-05, Train Loss = 34.764, Test Loss =
34.795	Iteration 29, LR: 6.25e-05, Train Loss = 34.699, Test Loss =
34.778	Iteration 30, LR: 3.125e-05, Train Loss = 34.240, Test Loss =
34.283	Iteration 31, LR: 3.125e-05, Train Loss = 34.187, Test Loss =
34.240	Iteration 32, LR: 3.125e-05, Train Loss = 34.183, Test Loss =
34.273	Iteration 33, LR: 3.125e-05, Train Loss = 34.084, Test Loss =
34.143	Iteration 34, LR: 3.125e-05, Train Loss = 34.048, Test Loss =
34.126	Iteration 35, LR: 3.125e-05, Train Loss = 34.067, Test Loss =
34.140	Iteration 36, LR: 3.125e-05, Train Loss = 33.938, Test Loss =
33.999	Iteration 37, LR: 3.125e-05, Train Loss = 33.972, Test Loss =
34.066	Iteration 38, LR: 3.125e-05, Train Loss = 33.906, Test Loss =
33.977	Iteration 39, LR: 3.125e-05, Train Loss = 33.930, Test Loss =
34.031	Iteration 40, LR: 1.5625e-05, Train Loss = 33.737, Test Loss =
33.846	

Done AutoEncoder Training

```
/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(  
/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-  
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(  

```

AE Acc: 90.40, PCA Acc: 89.98

Starting dimension reduction to 100 components

Starting PCA

Finished PCA

Starting AutoEncoder Training

411.124	Iteration 0, LR: 0.00025, Train Loss = 400.061, Test Loss =
66.640	Iteration 1, LR: 0.00025, Train Loss = 67.329, Test Loss =
56.282	Iteration 2, LR: 0.00025, Train Loss = 56.515, Test Loss =
45.705	Iteration 3, LR: 0.00025, Train Loss = 45.893, Test Loss =
76.851	Iteration 4, LR: 0.00025, Train Loss = 76.234, Test Loss =
40.700	Iteration 5, LR: 0.00025, Train Loss = 40.744, Test Loss =
43.148	Iteration 6, LR: 0.00025, Train Loss = 43.252, Test Loss =
38.480	Iteration 7, LR: 0.00025, Train Loss = 38.617, Test Loss =
40.016	Iteration 8, LR: 0.00025, Train Loss = 39.971, Test Loss =
37.274	Iteration 9, LR: 0.00025, Train Loss = 37.326, Test Loss =

32.562	Iteration 10, LR: 0.000125, Train Loss = 32.580, Test Loss =
32.687	Iteration 11, LR: 0.000125, Train Loss = 32.723, Test Loss =
31.100	Iteration 12, LR: 0.000125, Train Loss = 31.059, Test Loss =
29.885	Iteration 13, LR: 0.000125, Train Loss = 29.890, Test Loss =
33.379	Iteration 14, LR: 0.000125, Train Loss = 33.256, Test Loss =
28.517	Iteration 15, LR: 0.000125, Train Loss = 28.503, Test Loss =
28.110	Iteration 16, LR: 0.000125, Train Loss = 28.047, Test Loss =
27.527	Iteration 17, LR: 0.000125, Train Loss = 27.495, Test Loss =
27.194	Iteration 18, LR: 0.000125, Train Loss = 27.164, Test Loss =
27.188	Iteration 19, LR: 0.000125, Train Loss = 27.145, Test Loss =
25.715	Iteration 20, LR: 6.25e-05, Train Loss = 25.625, Test Loss =
25.715	Iteration 21, LR: 6.25e-05, Train Loss = 25.671, Test Loss =
25.677	Iteration 22, LR: 6.25e-05, Train Loss = 25.602, Test Loss =
25.429	Iteration 23, LR: 6.25e-05, Train Loss = 25.323, Test Loss =
25.385	Iteration 24, LR: 6.25e-05, Train Loss = 25.313, Test Loss =
25.301	Iteration 25, LR: 6.25e-05, Train Loss = 25.209, Test Loss =
25.142	Iteration 26, LR: 6.25e-05, Train Loss = 25.052, Test Loss =
24.948	Iteration 27, LR: 6.25e-05, Train Loss = 24.830, Test Loss =
24.994	Iteration 28, LR: 6.25e-05, Train Loss = 24.860, Test Loss =
24.886	Iteration 29, LR: 6.25e-05, Train Loss = 24.748, Test Loss =
24.390	Iteration 30, LR: 3.125e-05, Train Loss = 24.256, Test Loss =
24.339	Iteration 31, LR: 3.125e-05, Train Loss = 24.209, Test Loss =
24.193	Iteration 32, LR: 3.125e-05, Train Loss = 24.059, Test Loss =
24.247	Iteration 33, LR: 3.125e-05, Train Loss = 24.122, Test Loss =

```

Iteration 34, LR: 3.125e-05, Train Loss = 23.989, Test Loss =
24.119
Iteration 35, LR: 3.125e-05, Train Loss = 23.989, Test Loss =
24.117
Iteration 36, LR: 3.125e-05, Train Loss = 23.926, Test Loss =
24.066
Iteration 37, LR: 3.125e-05, Train Loss = 23.941, Test Loss =
24.077
Iteration 38, LR: 3.125e-05, Train Loss = 23.771, Test Loss =
23.893
Iteration 39, LR: 3.125e-05, Train Loss = 23.837, Test Loss =
23.979
Iteration 40, LR: 1.5625e-05, Train Loss = 23.562, Test Loss =
23.699

```

Done AutoEncoder Training

```

/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

AE Acc: 90.31, PCA Acc: 89.70

```

/home/apoorvsharma/anaconda3/envs/data598/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```


### 3.0.6 Output Analysis

Make a plot with  $d$  on the x-axis and the best test accuracy of the logistic regression model on the y-axis with the  $d$ -dimensional representations. The plot should have two lines, corresponding to PCA and autoencoders.



```
[31]: plt.plot(num_components_list, ae_accuracy, label="Auto-Encoder Accuracy")
plt.plot(num_components_list, pca_accuracy, label="PCA Accuracy")
plt.title('Dimension Reduction Analysis', fontsize=18)
plt.ylabel('Accuracy (%age)', fontsize=18)
plt.xlabel('Latent Features', fontsize=18)
plt.legend(fontsize=18)
```

```
[31]: <matplotlib.legend.Legend at 0x7f358027efa0>
```



Based on these observations, could you speculate why one of the two might be better or worse than the other?

PCA is essentially a linear transformation. However, AutoEncoders can map complex non-linear functions. As a result, for image data, auto-encoders perform better at reconstruction.